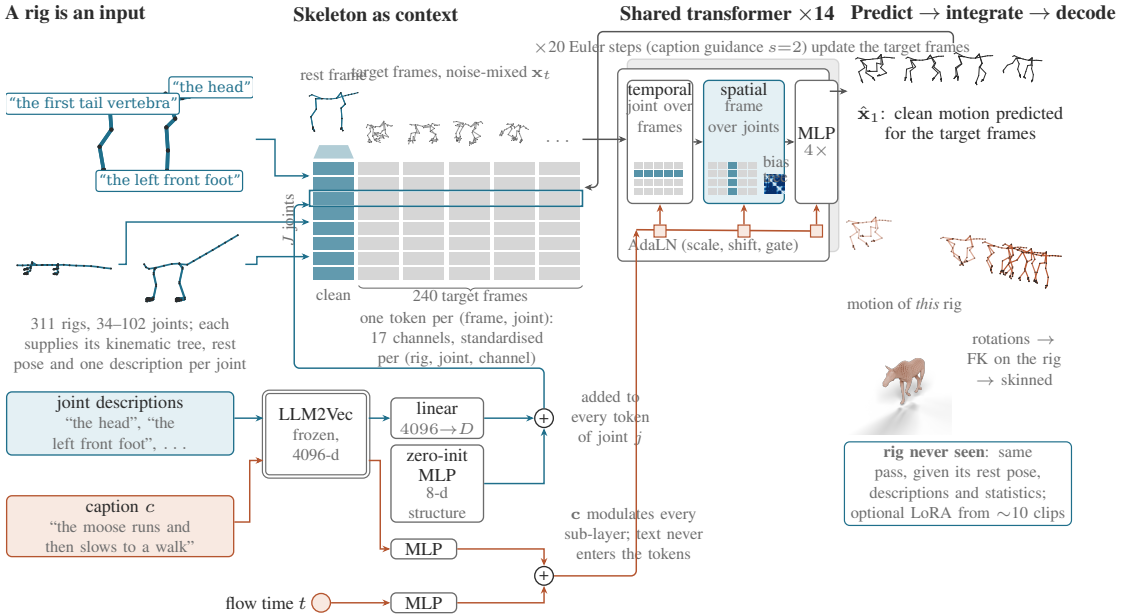

A rig is an input
Skeleton as context
Shared transformer ×14
Predict → integrate → decode
×20 Euler steps (caption guidance s=2) update the target frames
"the head"
"the first tail vertebra"
"the left front foot"
rest frame
target frames, noise-mixed $\mathbf{x}_t$
J joints
clean
240 target frames
311 rigs, 34–102 joints; each supplies its kinematic tree, rest pose and one description per joint
one token per (frame, joint): 17 channels, standardised per (rig, joint, channel)
temporal joint over frames
spatial frame over joints
bias
MLP 4×
AdaLN (scale, shift, gate)
$\hat{\mathbf{x}}_1$: clean motion predicted for the target frames
motion of *this* rig
rotations → FK on the rig → skinned
joint descriptions "the head", "the left front foot", . . .
LLM2Vec frozen, 4096-d
linear 4096→D
zero-init MLP 8-d structure
added to every token of joint j
caption c "the moose runs and then slows to a walk"
MLP
c modulates every sub-layer; text never enters the tokens
flow time t
MLP
rig never seen: same pass, given its rest pose, descriptions and statistics; optional LoRA from ∼10 clips